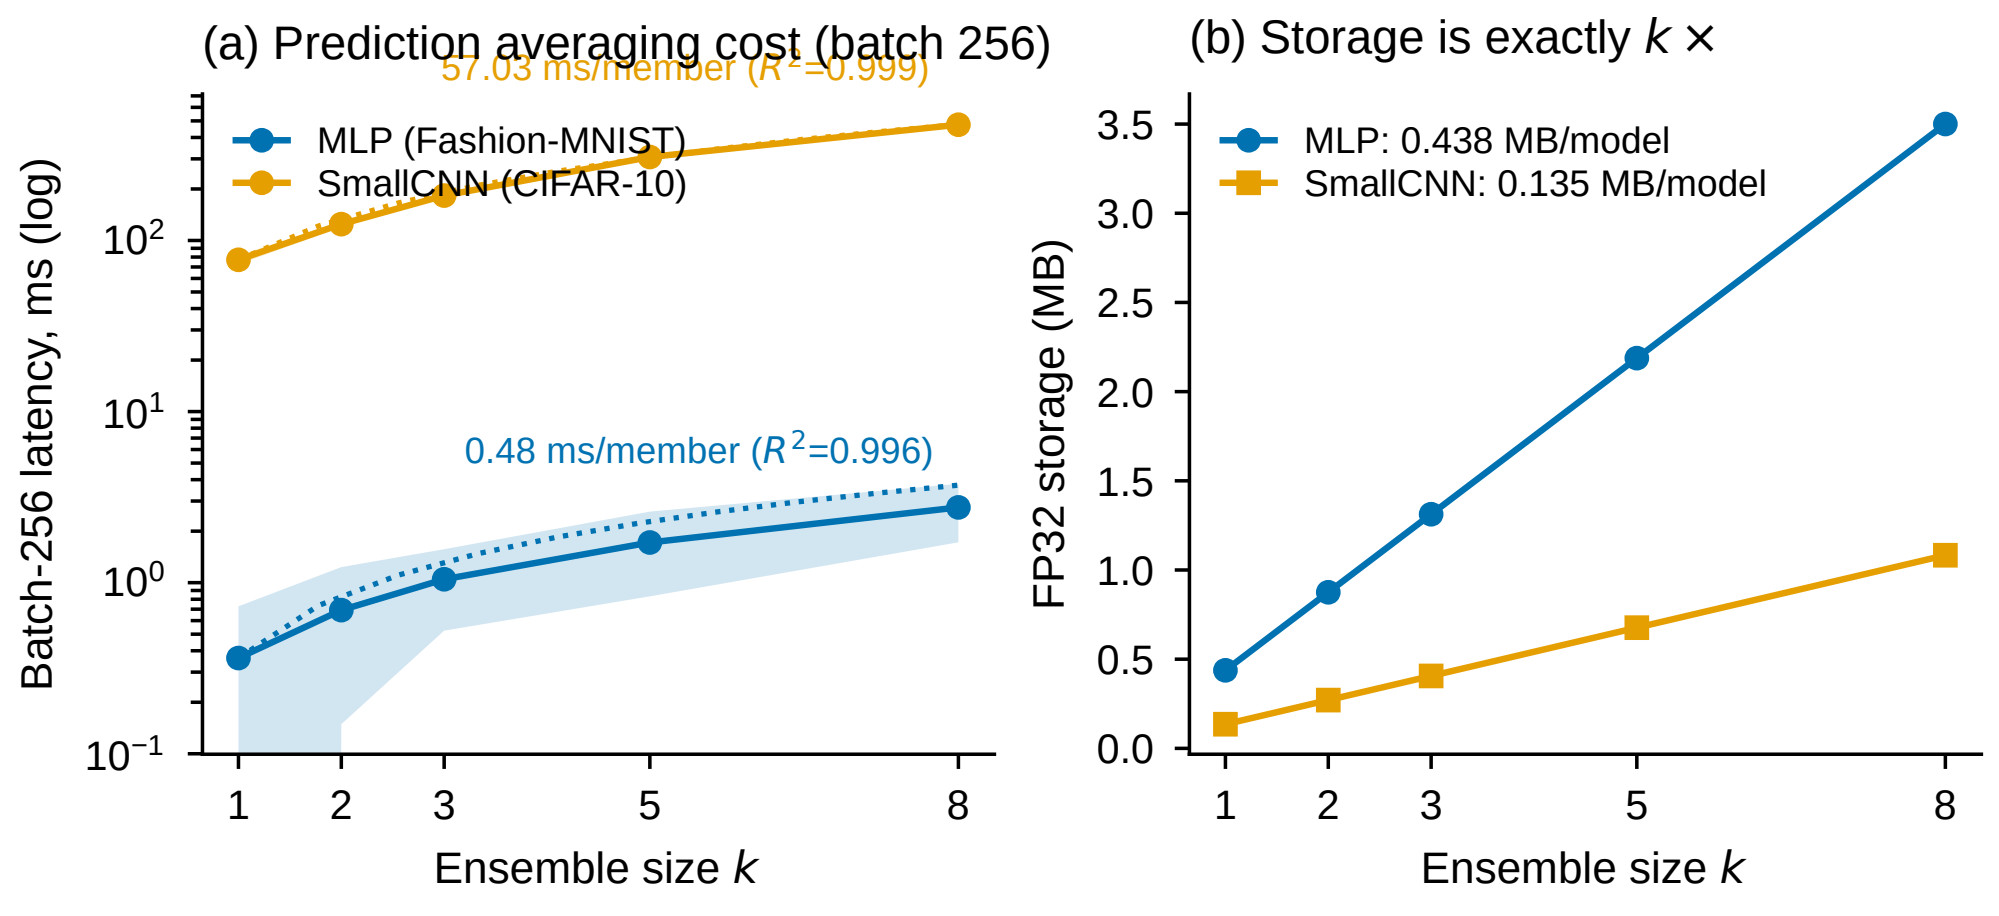
(a) Prediction averaging cost (batch 256)
57.03 ms/member (R^2=0.999)
MLP (Fashion-MNIST)
SmallCNN (CIFAR-10)
0.48 ms/member (R^2=0.996)
Batch-256 latency, ms (log)
10^2
10^1
10^0
10^{-1}
1
2
3
5
8
Ensemble size k
(b) Storage is exactly k ×
MLP: 0.438 MB/model
SmallCNN: 0.135 MB/model
FP32 storage (MB)
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
1
2
3
5
8
Ensemble size k